

Discussant Remarks

Andreas Kuepfer
TU Darmstadt

Annual Meeting of the European Political Science Society · 18–20 June 2026 · Belfast, UK

What Unites the Projects

Paper 1

Nalewajko & Spirig

Of Polish Heroism and Victimhood: How Wikipedia Captures Collective Memory

Paper 2

Gruber, Hertweck & Sältzer

Ghostcounting: A Bayesian Approach to Estimating Protester Counts and Media-Level Bias

*Both papers recover a **latent, contested truth** from **biased, multi-source text records**, treating bias as a parameter to model, not noise to average away.*

Paper 1: Of Polish Heroism and Victimhood

Of Polish Heroism and Victimhood: How Wikipedia Captures Collective Memory

Nalewajko & Spirig

Why This Is a Great Idea

Contribution

- A scalable, repeatable measure of collective memory content
- Survey / textbook / interview approaches don't travel across time or language

What Convinces Me

- Conceptual framework (role-to-group allocation) is portable to *any* contested past
- Wikipedia's dual role (record **and** influence)
- Built-in cross-language, cross-time comparison is “for free” once the pipeline exists

Comment 1: Who Gets to Write Memory?

The paper treats edits as a window onto the “average member’s beliefs.” But editors are not a random draw from the language community, and not every proposed edit survives.

- **Who edits?** Wikipedia editors are disproportionately male, technically literate, and skew toward specific demographics (Mandiberg 2020). Does this matter more for *contested* memory topics than for, say, math articles?
- **What gets accepted?** A framework built on *final article content* implicitly already filters through reverts, talk-page disputes, and admin moderation
- Idea: use **edit acceptance/reversion rates** by group-role content as a second outcome

Comment 2: Different Use Cases

Polish WWII memory is a best-case scenario: a topic with strong elite *and* popular consensus, mostly settled before Wikipedia existed.

- The paper needs at least one case to show the method detects *change*, not just stability
- Natural candidates with ongoing memory construction over the past years:
 - **Russia's invasion of Ukraine:** “liberation” vs. “invasion” framing, contested while being written
 - **Gaza war:** competing victim/perpetrator framings across language editions, high edit intensity, frequent protection
- These cases let you show the method's two failure/success modes side by side: consensus narrative (Poland) vs. live narrative war (Ukraine/Gaza)

Comment 3: Wikipedia's Gatekeeper Role Is Changing

- **LLM chatbots** increasingly mediate “what people know about the past”, trained on Wikipedia, but also competing sources (Grokopedia, etc.)
- If Wikipedia's influence is declining or filtered through an LLM layer, the paper's causal link (Wikipedia → public belief) needs updating
- Idea: how would this measurement strategy adapt if collective memory increasingly forms via LLM-mediated answers?

Minor Comments

- **Bots and non-human edits:** a sizeable share of Wikipedia edits are bot-driven (formatting, link fixes, vandalism reversion) and should be filtered before treating edit counts as proxies for “community belief”
- **Validate the LLM classification step itself**
- **Event-study validation:** use known memory-politics shocks (2018 Polish Memory Law, Jedwabne / Gross debate, 2015 PiS election) as natural experiments — does the group-role embedding move detectably around these dates? This would be strong external validation beyond matching the qualitative literature in levels
- **Reader-side, not just editor-side, validation:** pair pageview spikes with the group-role measure to ask whether *exposure* to a given narrative tracks survey-based collective memory measures
- **How memory changes:** Removals / replacements are stronger signals than additions; protection-level history + revert rates as a companion contestation measure
- **Cross-language & cross-editor:** Compare role-group associations across Polish, German, English, Russian Wikipedia; test whether bilingual editors export or independently construct narratives

Ghostcounting: A Bayesian Approach to Estimating Protester Counts and Media-Level Bias

Gruber, Hertweck & Sältzer

Why This Is a Great Idea

Contribution

- First fully automated, near-real-time protest database
- Disagreeing sources = the *signal* (bias becomes an estimated parameter)
- Local LLM: no API costs, fully reproducible

What Convinces Me

- The simulation shows ideological skew is identifiable *without* feeding the model ideology labels
- Modular pipeline (RSS -> scrape -> classify -> group -> estimate) is reusable for other event types, not just protests
- Public dashboard + open repo lowers the barrier for replication and reuse

Comment 1: Historical Depth and Visual Evidence

The pipeline is built to run forward from today. Two extensions would make it useful for studying the past, not just the present.

- **Historical backfill:** RSS only goes back to pipeline launch (retrospective crawling from newspaper archives?)
- **Photographic evidence:** Crowd photos (aerial shots, density estimates) as an independent signal alongside text-based counts?

Comment 2: Entity Matching and Data Access

Two related questions about what's identified inside an event and what's accessible outside the dashboard.

- **Entity matching:** Are protest groups (e.g., “Fridays for Future”) linked across events or treated as free text per event? Entity linking would let you track *movement* trajectories
- **Rehydration:** The dashboard exposes structured fields but not full article text — how would a researcher access the raw scraped content for follow-up analysis?

Comment 3: Where Did the Number Come From?

The Bayesian model currently treats “source” as the outlet. But the outlet is often just relaying someone else’s count.

- Many crowd-size figures in news text are themselves attributed to **police estimates, organizer claims, or news agencies (e.g., dpa)**, but the outlet didn’t generate the number
- Idea: prompt the (already-deployed) LLM classifier to additionally extract **attribution** (“according to police,” “organizers said,” no source given) as a structured field

Minor Comments

- **Validation against ground truth:** Benchmark $\hat{z}_i$ against an existing hand-coded protest dataset; separately validate event detection (Steps 1–2) from the Bayesian model (Steps 3–4)
- **Time-varying source bias:** Ownership changes or editor turnover can shift an outlet's reporting — a rolling-window extension ($\beta_{s,t} = \beta_{s,t-1} + \eta_{s,t}$) would let bias itself become a time series